

ENSF 593/594

Data Structures – AVL Trees

Mohammad Moshirpour

A series of horizontal lines in shades of green and white, located on the right side of the slide, extending from the left edge of the text area.

Outline

- AVL Tree
- Insertion to AVL Trees

Goal

- In this lecture we will learn about AVL trees which are self balancing trees and make searches and insertions more efficient.

The Need for Balanced Trees

- Searches and insertions are most efficient when a binary tree is well balanced.
- Binary trees may become unbalanced after insertions and deletions.
 - In the worst case, the tree degenerates into a linked list.
- There are several variants of ordered binary trees that remain well balanced after insertions and deletions.
 - AVL trees are one example.


Balance of a Node

- **Definition:** is the height of the right subtree minus the height of the left subtree:





$$\text{balance}(n) = \text{rightHeight}(n) - \text{leftHeight}(n)$$

where n is some node in the tree

□ E.g.




Balance of a Node (cont'd)






AVL Trees

- Named after their inventors: G. M. Adel'son-Vel'skii and E. M. Landis
 - Developed in 1962
- Definition: is an ordered binary tree where every node has a balance of -1, or 0, or +1
 - E.g.





AVL Trees (cont'd)



AVL Trees (cont'd)

- Note that the difference between the subtrees can never exceed 1
- Examples of non-AVL trees



Node Structure

- Must add a *balance* field to the Node class used for binary search trees:

```
public class Node {  
    private int data;  
    private Node parent, left, right;  
    private int balance;  
  
    . . .  
}
```

Insertion into an AVL Tree

- General procedure:
 - 1) Insert the node into the tree, following the rules for a regular binary search tree
 - 2) If necessary, adjust the shape of the tree so that it conforms to the rules of an AVL tree
 - Involves doing a single or double *rotation*
 - 3) Update the balance fields for all nodes affected by the steps above

Insertion (cont'd)

- *Pivot Node*
 - **Definition:** is the ancestor node closest to the inserted node that is *not* in balance
 - i.e. is not 0
 - It is possible there may be no pivot when doing an insertion
 - One adjusts the AVL tree and updates the balances according to the nature of the pivot and where the insertion is done

Insertion (cont'd)

- There are 3 possible cases when doing an insertion:
 - 1) There is no pivot
 - 2) The pivot exists, and you add to the shorter subtree
 - 3) The pivot exists, and you add to the longer subtree

Insertion (cont'd)

- **Case 1: There is no pivot**
 - Essentially, you are adding to a subtree with all 0 balances
 - You change the balances for all ancestor nodes by ± 1
 - The shape of the tree is *not* adjusted after the insertion
 - But the balances must be updated


Insertion (cont'd)

□ Procedure:

- Insert the node into its proper place in the tree
- Adjust the balances for all nodes from the inserted node up to the root node (i.e. all nodes on the *search path*)
 - The inserted node is given a balance of 0
 - For the other nodes:
 - If inserted node < node value, decrement balance
 - If inserted node > node value, increment balance


Insertion (cont'd)

- E.g.
 - Original tree




Insertion (cont'd)

- Insert E:




Insertion (cont'd)

- Update balances:




Insertion (cont'd)

- E.g.
 - Original tree




Insertion (cont'd)

- Insert H



Insertion (cont'd)

- Update balances



Insertion (cont'd)

- **Case 2: A pivot exists, and a node is added to the shorter subtree**
 - Essentially, you are adding to a shorter subtree to bring it into better balance
 - The shape of the tree is *not* adjusted after the insertion
 - But the balances must be updated

Insertion (cont'd)

- You must be able to tell if you are adding to the shorter subtree (to distinguish from Case 3); you are if:
 - $\text{Pivot} == +1$ and $\text{inserted node} < \text{pivot node}$, or
 - $\text{Pivot} == -1$ and $\text{inserted node} > \text{pivot node}$


Insertion (cont'd)

□ Procedure:

- Insert the node into its proper place in the tree
- Adjust the balances for all nodes from the inserted node up to and including the *pivot* node
 - The inserted node is given a balance of 0
 - For the other nodes:
 - If inserted node < node value, decrement balance
 - If inserted node > node value, increment balance
 - Note that balances do not change above the pivot node


Insertion (cont'd)

- E.g.
 - Original tree




Insertion (cont'd)

- Insert A, identify pivot




Insertion (cont'd)

- Adjust balances




Insertion (cont'd)

- E.g.
 - Original tree




Insertion (cont'd)

- Add F, identify pivot



Insertion (cont'd)

- Adjust balances



Insertion (cont'd)

- **Case 3: A pivot exists, and you add to the longer subtree**
 - Essentially, you are putting the tree into worse balance
 - The pivot's balance changes to ± 2
 - The shape of the tree *must* be adjusted after doing the insertion

Insertion (cont'd)

- You must be able to tell if you are adding to the longer subtree (to distinguish from Case 2); you are if:
 - Pivot == +1 and inserted node > pivot node, or
 - Pivot == -1 and inserted node < pivot node
- Case 3 breaks down into 2 subcases:
 - You add to the “outside subtree” of the “son” of the pivot on the search path
 - You add to the “inside subtree” of the “son” of the pivot on the search path

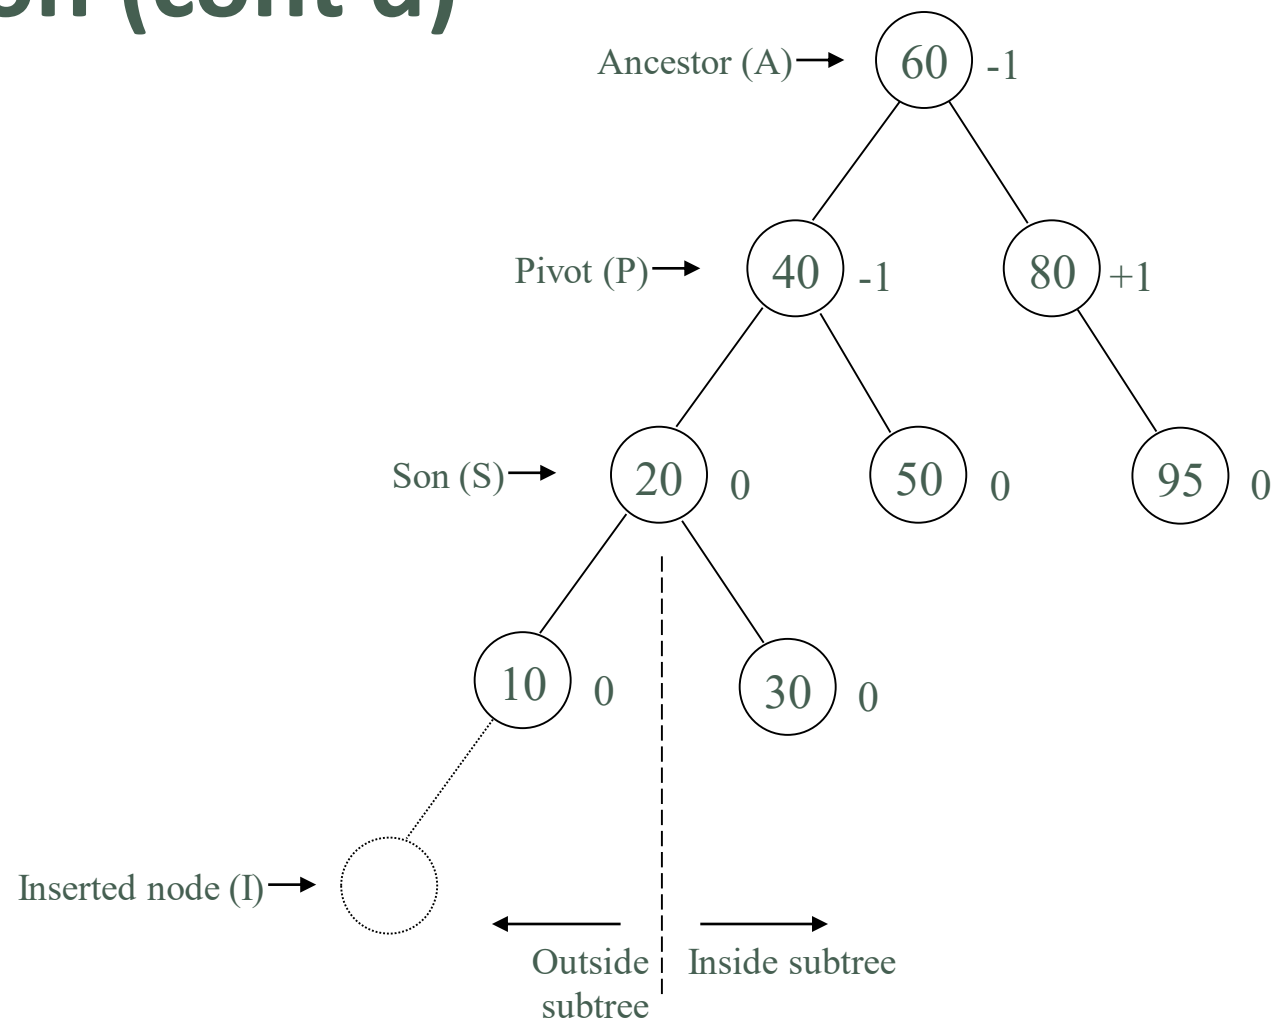
Insertion (cont'd)

- **Terminology:**

- *Ancestor node*: the parent node of the pivot node
- *Son node*: the child node of the pivot node, on the path from the pivot to the inserted node
- *Outside subtree*:
 - The left subtree of the son, if the pivot is negative
 - The right subtree of the son, if the pivot is positive

Insertion (cont'd)

□ E.g.



Insertion (cont'd)

- **Case 3a: adding a node to the outside subtree**
 - Procedure:
 - 1) Insert the node into its proper place in the tree
 - 2) Adjust the shape of the tree by doing a *single rotation*. 2 cases:
 - a) Do a *right rotation* if the outside subtree is on the left (the pivot is negative)
 - b) Do a *left rotation* if the outside subtree is on the right (the pivot is positive)

Insertion (cont'd)

3 pointers must be changed:


- 1) If $\text{pivot} < \text{ancestor}$, then ancestor's *left* child pointer is set to the son node, otherwise set the *right* child pointer
- 2) 2 cases:
 - a) Right rotation: pivot's left child pointer set to the right child of the son node
 - b) Left rotation: pivot's right child pointer is set to the left child of the son node
- 3) 2 cases:
 - a) Right rotation: son's right child pointer set to the pivot node
 - b) Left rotation: son's left child pointer set to the pivot node

Insertion (cont'd)

- 3) Adjust balances of affected nodes:
 - a) Set pivot and inserted node to 0
 - b) Adjust the balances for all nodes above the inserted node, up to the child of the son node
 - If inserted node $<$ node value, decrement balance
 - If inserted node $>$ node value, increment balance


Insertion (cont'd)

- E.g.
 - Original tree




Insertion (cont'd)

- Add 5, identify nodes and subtrees





Insertion (cont'd)


- Do right rotation (3 substeps)



Insertion (cont'd)




Insertion (cont'd)




Insertion (cont'd)

- Adjust balances




Insertion (cont'd)

- E.g.
 - Original tree




Insertion (cont'd)

- Add Q, identify nodes and subtrees




Insertion (cont'd)

- Do left rotation



- Adjust balances




Insertion (cont'd)

- **Case 3b: adding a node to the inside subtree**
 - *Grandson node*: the child node of the son node, on the path from the pivot to the inserted node
 - To adjust the tree after an insertion, a *double rotation* is performed; consists of:
 - A right rotation at one node, followed by a left rotation at another node (RL rotation), or
 - The inverse (LR rotation)


Insertion (cont'd)

- Example: LR double rotation




Insertion (cont'd)

- Left rotation at son



Insertion (cont'd)

- Right rotation at pivot



Insertion (cont'd)

- **Procedure:**
 - 1) Insert the node into its proper place in the tree
 - 2) Adjust the shape of the tree by doing a double rotation; 2 cases:
 - a) RL rotation if the pivot is positive
 - b) LR rotation if the pivot is negative

Insertion (cont'd)

- RL rotation:
 - 1) Right rotation through son
 - a) Set pivot's right child pointer to grandson node
 - b) Set son's left child pointer to grandson's right subtree (if it exists)
 - c) Set grandson's right child pointer to son node
 - 2) Left rotation through pivot
 - a) If there is no ancestor, set the root pointer to grandson; if $\text{pivot} > \text{ancestor}$, set the ancestor's right child pointer to the grandson, else set the left child pointer

Insertion (cont'd)

- b) Set pivot's right child pointer to grandon's left subtree (if it exists)
 - c) Set grandson's left child pointer to pivot
- LR rotation is symmetrical to the RL rotation

Insertion (cont'd)

3) Adjust balances of affected nodes

a) Set inserted node to 0

b) RL rotation:

- If inserted node $>$ grandson, set pivot to -1
- Else set pivot to 0, son to +1


c) LR rotation is symmetrical to the above

d) Adjust balances for all nodes above the inserted node up to the child of the son or pivot

- If inserted node $<$ node value, decrement balance
- If inserted node $>$ node value, increment balance


Insertion (cont'd)

- E.g.
 - Original tree





Insertion (cont'd)

- Add E, identify nodes and subtree





Insertion (cont'd)


- Right rotation at son (3 substeps)



Insertion (cont'd)





Insertion (cont'd)




Insertion (cont'd)

- Left rotation at pivot (3 substeps)




Insertion (cont'd)



Insertion (cont'd)

- Adjust balances



Summary

- AVL Trees: the difference between the subtrees can never exceed 1.
- Insertion:
 - There is no pivot
 - The pivot exists, and add to the shorter subtree
 - The pivot exists, and add to the longer subtree
 - adding a node to the outside subtree
 - adding a node to the inside subtree

Review Questions

- How can you calculate the Balance of a Node?
- What is an AVL tree?
- What is the difference between nodes in a binary and AVL trees?
- What is the general procedure of inserting nodes to an AVL tree?
- What is a Pivot Node?
- What should you do when there is no pivot while adding a node to an AVL?
- What should you do when a pivot exists, and a node is added to the shorter subtree while adding a node to an AVL?
- What should you do when A pivot exists, and you add to the longer subtree while adding a node to an AVL?



Any questions?